

# Midterm Project

Ethan Alteza, Barron Bronson, Holden Ellis, Charlotte Huang

2026-02-08

## Write Up

<https://docs.google.com/document/d/1HZTEfuxbhKEsMSGOfqDvulxrU5hy9rIVp31Wf9el38g/edit?usp=sharing>

```
source("Ziggurat.R")
source("BoxMuller.R") # if anyone wants to compare, much faster because of how R runs code

# build the table for the Ziggurat
x <- zigtable(function(x) {1/sqrt(2*pi) * exp(-x**2/2)}, function(y) {sqrt(-2*log(y*sqrt(2*pi)))}, 8)$x
y <- dnorm(x)
r_val <- tail(x, 1)

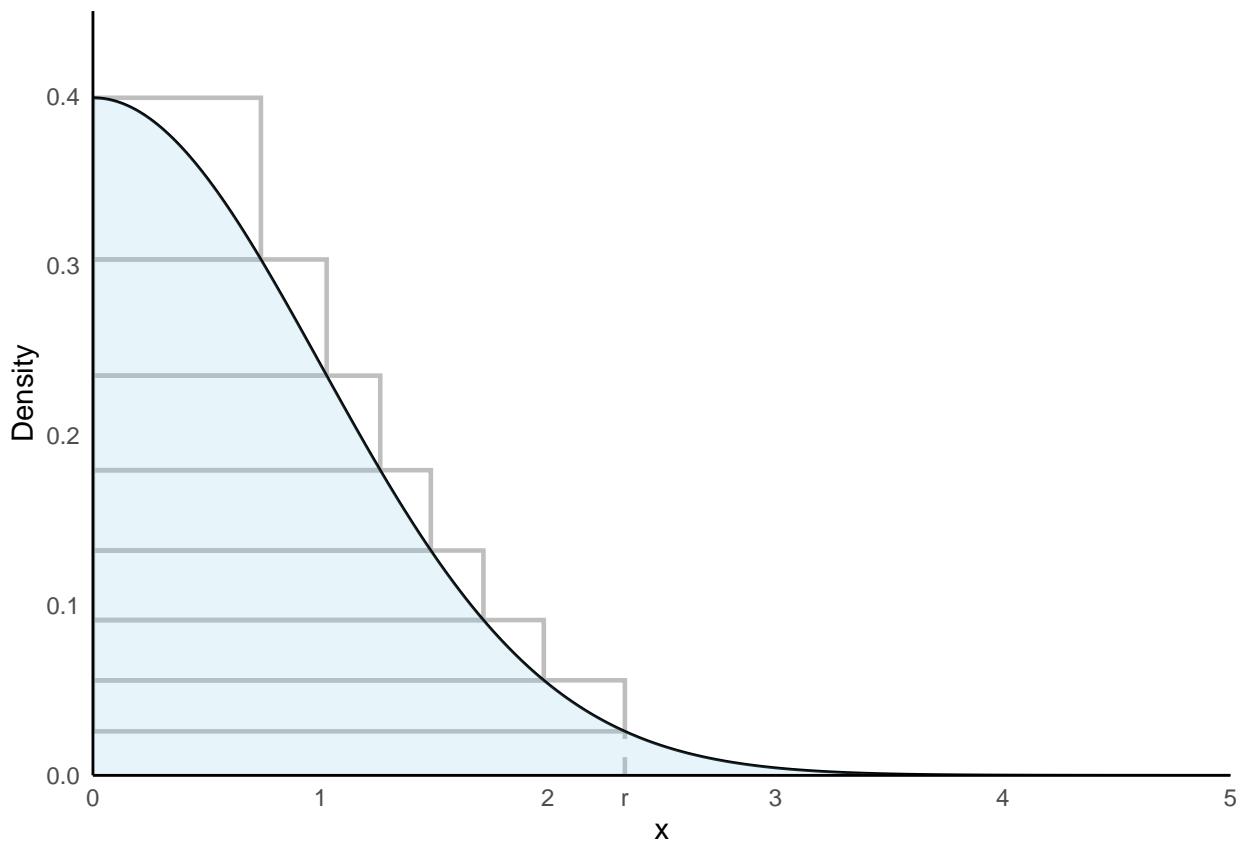
# Dataframe to draw rectangles
df_rect <- data.frame(
  xmin = 0,
  xmax = x[-1],           # x[2]...x[8]
  ymin = y[-1],           # y[2]...y[8]
  ymax = y[-length(y)]   # y[1]...y[7]
)

ggplot() +
  # draw rectangles
  geom_rect(data = df_rect,
            aes(xmin = xmin, xmax = xmax, ymin = ymin, ymax = ymax),
            fill = NA, color = "grey", linewidth = 0.8) +
  
  # dotted line at x=r
  geom_segment(aes(x = r_val, xend = r_val, y = 0, yend = tail(y, 1)),
               linetype = "dashed", color = "grey", linewidth = 0.8) +
  
  # draw normal distribution
  stat_function(fun = dnorm, geom = "line", color = "black", n = 1000) +
  # highlight area under the curve
  stat_function(fun = dnorm, geom = "area", fill = "skyblue", alpha = 0.2, xlim = c(0, 5)) +
  
  # style axes
  scale_x_continuous(
    limits = c(0, 5),
    expand = c(0, 0),
    breaks = sort(c(0:5, r_val)),
    labels = function(b) {
      # Use a small epsilon for float comparison
    }
  )
```

```

        ifelse(abs(b - r_val) < 1e-6, "r", b)
    }
) +
scale_y_continuous(limits = c(0, 0.45), expand = c(0, 0)) +
# labels and plot styles
labs(x = "x", y = "Density") +
theme_minimal() +
theme(
  panel.grid = element_blank(),
  axis.line = element_line(color = "black")
)

```



```

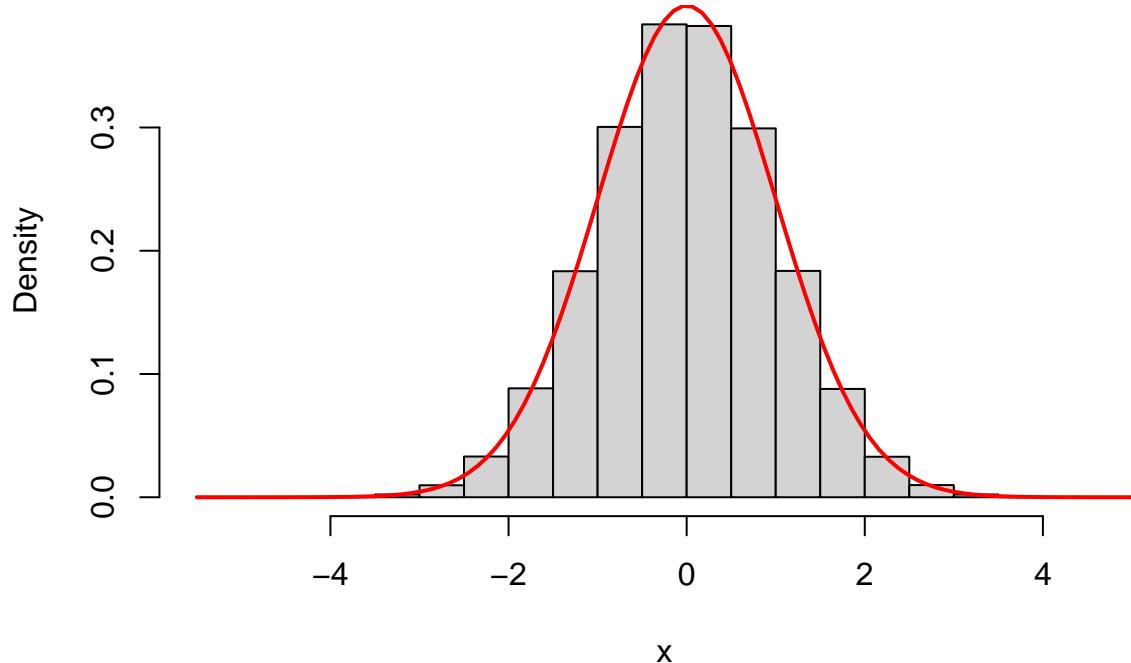
zigtatable(function(x) { exp(-x) }, function(y) { -log(y) }, 8)

## $x
## [1] 0.0000000 0.4338930 0.7888599 1.1387148 1.5163812 1.9560291 2.5166676
## [8] 3.3489677
##
## $v
## [1] 0.1527383

x <- rnормzig(1000000)
hist(x, freq=FALSE)
curve(dnorm(x), add=TRUE, col="red", lwd=2) # looks good

```

## Histogram of x



```
set.seed(777)

#testing: alpha = 0.01
samples1 <- rnormzig(4999)
print("4999 Samples")

## [1] "4999 Samples"
shapiro.test(samples1) # no evidence against normal; p-value > alpha = 0.01

##
## Shapiro-Wilk normality test
##
## data: samples1
## W = 0.99954, p-value = 0.2852
ks.test(samples1, "pnorm", mean=mean(samples1), sd=sd(samples1)) # same

##
## Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: samples1
## D = 0.0083523, p-value = 0.8766
## alternative hypothesis: two-sided
print("1,000,000 ziggurat samples")

## [1] "1,000,000 ziggurat samples"
```

```

# testing 1 mil samples...
ks.test(x, "pnorm", mean=mean(x), sd=sd(x)) # passed

##
##  Asymptotic one-sample Kolmogorov-Smirnov test
##
## data:  x
## D = 0.00052427, p-value = 0.9463
## alternative hypothesis: two-sided
mean(x) # close to 0

## [1] -0.0007271692
sd(x) # close to 1

## [1] 0.9998402
print("1,000,000 rnorm samples")

## [1] "1,000,000 rnorm samples"
# additional test:
x2 <- rnorm(1000000)
ks.test(x2, "pnorm", mean=mean(x2), sd=sd(x2)) # passed

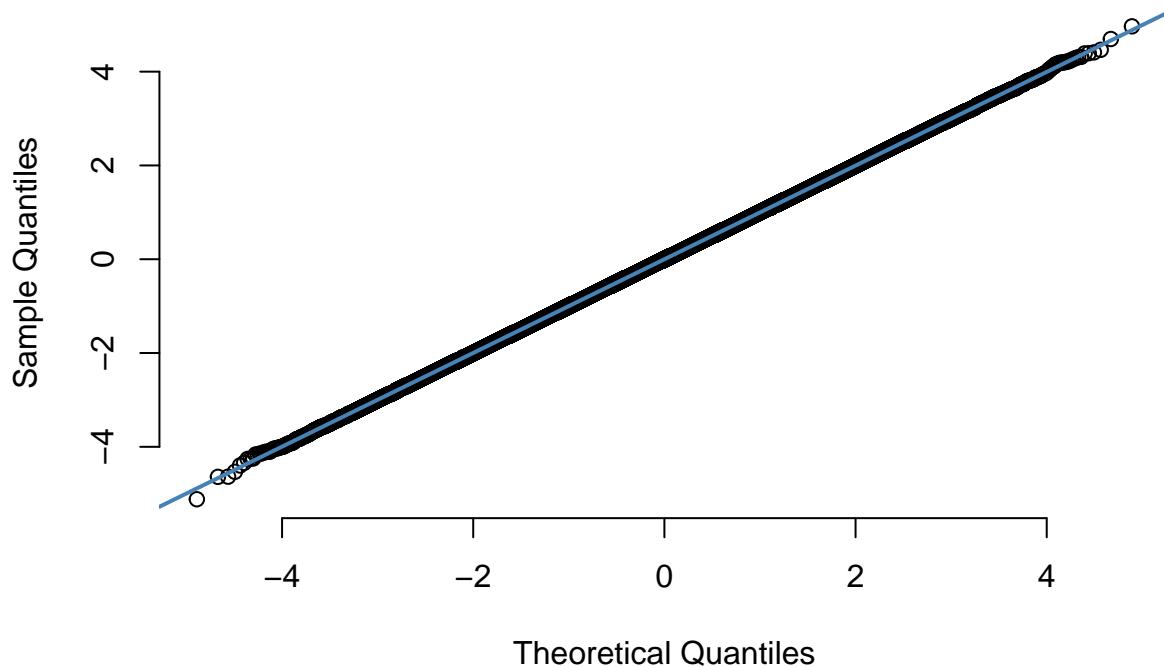
##
##  Asymptotic one-sample Kolmogorov-Smirnov test
##
## data:  x2
## D = 0.00069352, p-value = 0.722
## alternative hypothesis: two-sided
mean(x2) # close to 0

## [1] 0.000669734
sd(x2) # close to 1

## [1] 1.001078
# QQ plot
qqnorm(x, pch = 1, frame = FALSE)
qqline(x, col = "steelblue", lwd = 2) # basically normal

```

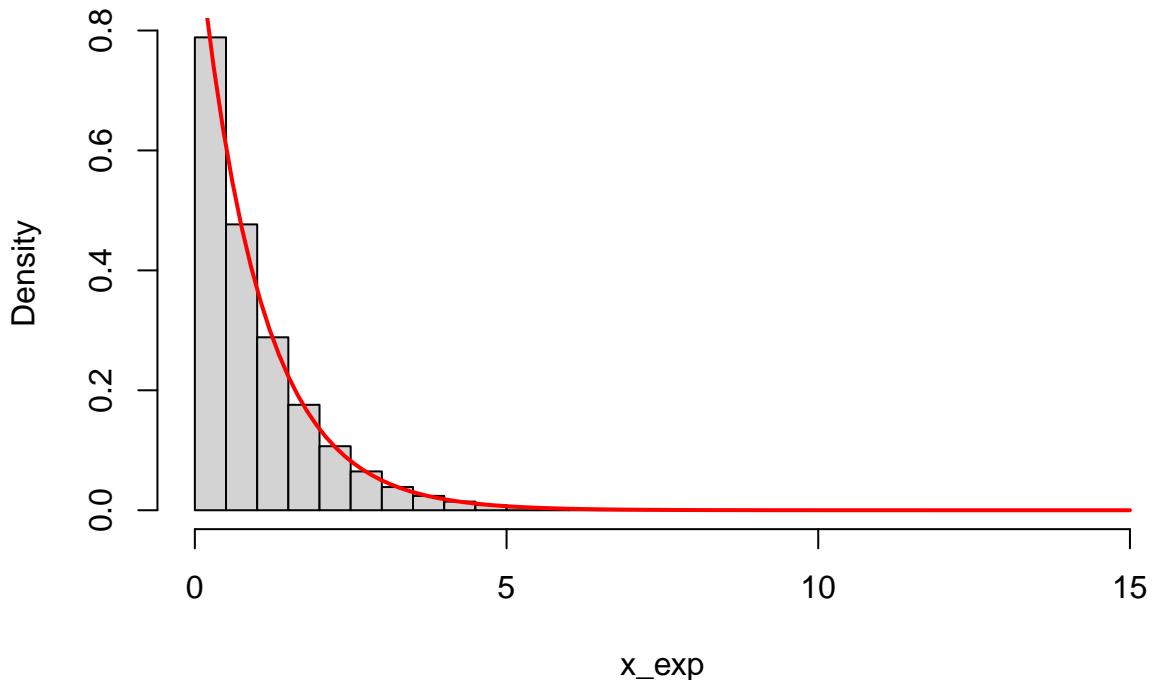
## Normal Q-Q Plot



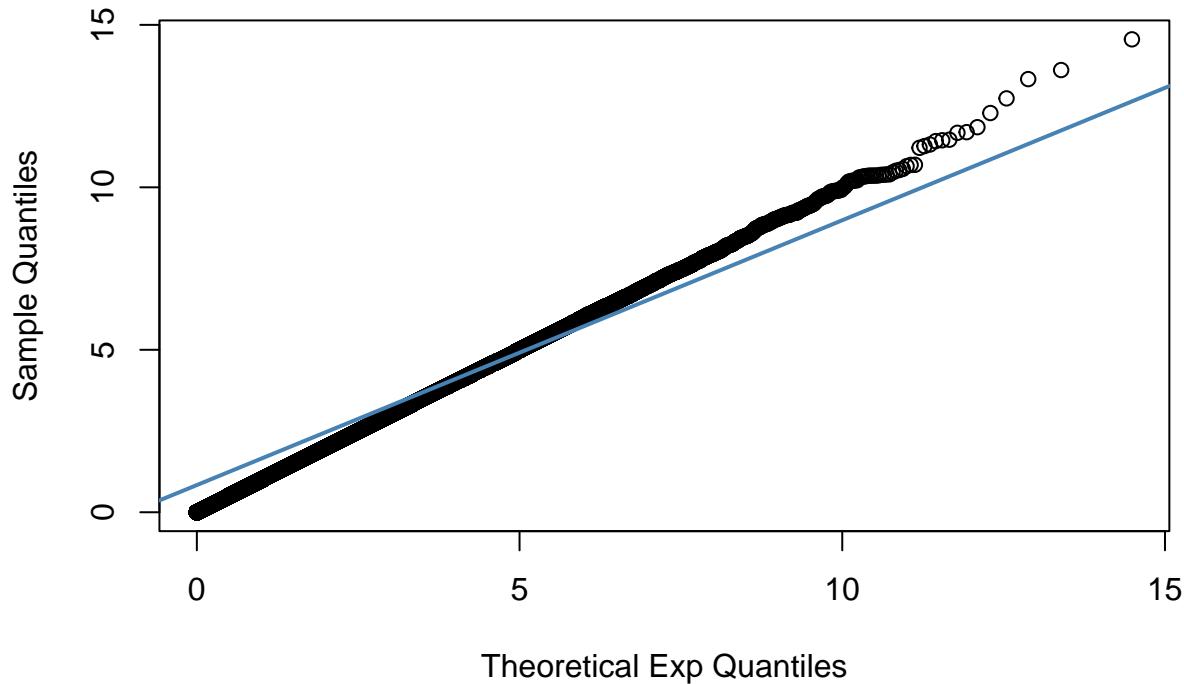
Assume we generate 100000 with Ziggurat method and they're actually normal, using rnorm() for now

```
set.seed(777)
x_exp <- rexpzig(1000000)
hist(x_exp, freq = FALSE)
curve(dexp(x), add=TRUE, col="red", lwd=2) # looks good
```

## Histogram of x\_exp



```
set.seed(777)
qqplot(qexp(ppoints(length(x_exp)), rate = 1/mean(x_exp)),
       sort(x_exp),
       xlab = "Theoretical Exp Quantiles",
       ylab = "Sample Quantiles")
qqline(x_exp, col = "steelblue", lwd = 2) # basically normal
```



```

print("1,000,000 Ziggurat samples")

## [1] "1,000,000 Ziggurat samples"
mean(x_exp) #Should be close to 1

## [1] 0.9985969
sd(x_exp) #Should be close to 1

## [1] 0.9979412
ks.test(x_exp, "pexp", rate = mean(x_exp)) #large p-value -> fail to reject exponential distribution

##
##  Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: x_exp
## D = 0.0014393, p-value = 0.03175
## alternative hypothesis: two-sided
print("1,000,000 rnorm samples")

## [1] "1,000,000 rnorm samples"
# additional test:
x2_exp <- rexp(1000000)
ks.test(x2_exp, "pexp", rate = mean(x2_exp)) # passed

## Warning in ks.test.default(x2_exp, "pexp", rate = mean(x2_exp)): ties should

```

```

## not be present for the one-sample Kolmogorov-Smirnov test
##
## Asymptotic one-sample Kolmogorov-Smirnov test
##
## data: x2_exp
## D = 0.00064725, p-value = 0.7963
## alternative hypothesis: two-sided
mean(x2_exp) # basically 1

## [1] 0.9999385
sd(x2_exp) # 1

## [1] 1.001024

```

A chi-squared goodness of fit test (Knuth 1997) is used to check the distribution of the random variables generated. This test quantizes the horizontal axis of the probability density function into k bins and derives a single value as a quality metric from the determined actual and expected number of samples in each bin.

$$\chi^2_{k-1} = \sum_{i=1}^k \frac{(Y_i - tp_i)^2}{tp_i}$$

Where t = number of observation,  $p_i$  = probability that each observation falls into the category i  $Y_i$  = the number of observation that actually do fall into the category i.

```

set.seed(999)

n = 100000
samples <- rnorm(n)

chi_squared_test <- function(x, k){
  t <- length(x) # number of observations
  bins <- seq(min(x), 7, length.out = k+1) #200 bins spaced uniformly over data edge of min and max

  Yi <- hist(x, breaks = bins, plot = FALSE)$counts # observed counts
  # expected probabilites (p_i) of standard normal
  p_i <- numeric(k)
  for (i in 1:k) {
    p_i[i] <- pnorm(bins[i+1]) - pnorm(bins[i])
  }
  # using computing chi square using equation from the paper
  chi_sq <- sum((Yi - t * p_i)^2 / (t * p_i))
  df <- k - 1
  return(data.frame(chi_square = chi_sq,
                    df = df))
}

set.seed(777)

k <- 200
results <- chi_squared_test(x, k)

```

```

results2 <- chi_squared_test(x2, k)

results$Method <- "Ziggurat"
results2$Method <- "Rnorm"

combined <- rbind(results, results2)
combined <- combined[, c("Method", setdiff(names(combined), "Method"))]
kable(combined, caption = "Chi-squared value: Ziggurat vs Rnorm")

```

Table 1: Chi-squared value: Ziggurat vs Rnorm

Method	chi_square	df
Ziggurat	176.7659	199
Rnorm	178.3372	199

The critical value for a chi-square test with 199 degrees of freedom at 95% confidence level ( $\alpha = 0.05$ ) is 232.912. (using a look up table on <https://www.medcalc.org/en/manual/chi-square-table.php>)